



AI-Powered Intelligent Job Search and Career System

Data Science and AI Project

Milestone 2 Report

Group 4

Prepared By:

Gaurav Kumar (22f1001105)
Dev Gupta (22f2000888)
Abhinav Ohri (24f1002064)
Pranav N (22f2000117)
Praveena N (22f3001454)



Index

SN	Particulars	Page No
1	Introduction	2
2	Dataset Identification	2
3	Dataset Description	4
4	Data Governance	7
5	Exploratory Data Analysis	9
6	Data Preprocessing	13
7	Dataset Integration	17
8	Data Augmentation	18
9	Dataset Splitting	19
10	Challenges Encountered	19
11	Deliverable produced	20
12	Summary and Next Steps	20
13	Team Consent	21

1. Introduction

The project, **AI-Powered Intelligent Job Search and Career System** aim to support candidates during the job search process by connecting resume parsing, candidate profile creation, career recommendation, and job discovery/matching into one structured workflow. In Milestone 1, the project scope was finalized around three main modules: **Resume Parsing**, **Career Recommendation**, and **Job Discovery and Job Matching**. The system first extracts structured information from a resume, then uses that profile to recommend suitable career roles, and finally ranks relevant jobs based on profile fit.

The objective of **Milestone 2** is to identify, understand, clean, preprocess, and prepare datasets for these modules.

2. Dataset Identification

2.1 Dataset List

Module	Dataset Name	Source	Status	Purpose
Resume Parsing	Resume Data (PDF)	Kaggle — Hadikp/resume-data-pdf	Publicly accessible dataset; exact licence to be recorded from the dataset page	Provide real resume documents for document-quality analysis, multimodal parsing, structured field extraction, and gold-set evaluation
Career Recommendation	ESCO Dataset v1.2.1 — English CSV classification	ESCO Official Portal	Public dataset	Map candidate skills and experience to suitable career roles
Job Discovery and Matching	Candidate Job Role Dataset and LinkedIn Job Postings Dataset	Kaggle	Public datasets	Provide candidate profiles and real-world job postings for semantic matching, feature engineering, ranking, and evaluation

2.2 Resume Parsing Dataset

The Resume Parsing module uses the Resume Data (PDF) dataset published on Kaggle. The dataset contains 8,905 resume documents organised into occupation-related folders. Unlike text-only resume datasets, this dataset contains complete resume documents and therefore allows the project to evaluate the actual document parsing workflow, including layout, image quality, and text readability.

The dataset was selected because it provides broad occupational coverage and contains 8,905 resumes across 97 raw folder labels. After correcting spelling and naming variations, these folders were mapped to 43 canonical job categories.

An important characteristic of the dataset is that all resumes are image-only PDFs. Each PDF contains one raster image and no extractable text layer. Therefore, ordinary PDF text extraction cannot be used as the main parsing method. The parser must use a multimodal vision model or an OCR-based fallback.

2.3 Career Recommendation Dataset

The ESCO dataset was selected because it provides a structured taxonomy of occupations, skills, knowledge, and abilities. This matches the design of the Career Recommendation module, where candidate profiles are compared with ESCO occupation vectors using embeddings and cosine similarity.

The ESCO portal provides downloadable data in formats such as CSV, JSON-LD, ODS, RDF, TTL, and XML. The project uses the English CSV classification dataset because it is easy to process using Python and suitable for preparing embedding-ready text fields.

2.4 Job Discovery and Matching Dataset

The Job Discovery and Matching module uses two complementary datasets. The Candidate Job Role Dataset contains structured candidate profiles with skills, qualifications, experience levels, and target job roles. The LinkedIn Job Postings Dataset contains real-world job titles, descriptions, companies, locations, required experience levels, salary-related fields, and employment types. Together, the datasets represent the two sides of the matching task: the candidate profile acts as a query and the job posting acts as a document.

2.5 Alternatives Considered

Alternative	Reason Considered	Reason Not Used as Main Dataset
O*NET	Well-known occupation and skill database	ESCO provides direct occupation-skill relation files
Snehaan Bhawal Resume Dataset	Contains more than 2,400 resume PDFs and category labels	Provides fewer job categories and less document diversity
ResumeAtlas	Larger dataset with 13,000+ resume records	Mainly provides OCR-extracted text and does not preserve the original resume document layout

3. Dataset Description

3.1 Resume Parsing Dataset

Item	Details
Dataset name	Resume Data (PDF)
Source	Kaggle — hadikp/resume-data-pdf
Raw records	8,905 resume PDFs
Format	Single-page, image-only PDF documents
Raw category folders	97
Canonical categories	43
Corrupt or unreadable files	0
Files with extractable PDF text	0
Exact duplicate files	0
Perceptual groups	6,611
Gold evaluation PDFs	86
Development set	34
Test set	52
Input feature	Rendered resume-page image
Weak label	Canonical resume category
Target fields	Contact, skills, education, experience, projects, certifications, and job titles

3.1.2 Resume Parsing target schema

Target group	Extracted fields
contact	Name, email, phone, location, links
skills	Skills, software tools, technologies, and competencies
education	Degree, field, institution, start year, and end year
experience	Job title, company, location, dates, current-role flag, and description
projects	Project name, description, and technologies
certifications	Certification name, issuer, and year

job_titles	Normalised previous or current role titles
------------	--

All 8,905 resumes are single-page PDFs containing an embedded image. No document contains a usable text layer. The median embedded image has a resolution of approximately $962 \times 1,600$ pixels, while the median PDF size is approximately 231 KB.

OCR was used only to estimate document length for EDA. On a reproducible sample of 200 resumes, the median resume contained approximately:

- 374 words;
- 2,699 characters;
- approximately 500 text tokens.

The 95th percentile was approximately 1,655 tokens, so most resumes can be processed in one model request without document chunking.

3.2 Career Recommendation Dataset

3.2.1 Raw Files Used

File	Purpose
occupations_en.csv	Contains ESCO occupation titles, alternative titles, descriptions, and occupation URIs
skills_en.csv	Contains ESCO skill titles, alternative skill names, descriptions, skill types, and skill URIs
occupationSkillRelations_en.csv	Connects occupations with related skills
ISCOGroups_en.csv	Provides broader occupation grouping, kept for reference
skillGroups_en.csv	Provides broader skill grouping, kept for reference

3.2.2 Processed Dataset Summary

Item	Value
Clean occupations	3,039
Clean skills	13,939
Clean occupation-skill relations	126,051
Final prepared career role records	3,039
Occupations with zero skills	0
Average skills per occupation	41.48
Unused skills	464

3.2.3 Final Career Recommendation Schema

Column	Description
occupation_id	Stable ESCO occupation ID extracted from concept URI
conceptUri	Original ESCO occupation URI
occupation_title	Main ESCO occupation title
alt_occupation_title	Alternative occupation titles
occ_description	ESCO occupation description
num_skills	Number of skills linked to the occupation
skills_list	Full list of linked ESCO skills
skills_for_embedding_text	Selected skills used in embedding text
career_role_context	Final combined text used for embedding generation

3.3 Job Discovery and Job Matching data

Metric	Candidate Job Role Dataset	LinkedIn Job Postings Dataset
Raw records	1,000	123,849
Raw features	5	31
Features retained	5	14
Data format	CSV	CSV
Native target	job_role	None
Ranking target	Pair-level relevance label	Pair-level relevance label

3.3.2 Candidate Job Role Features:

Feature	Description
candidate_id	Unique candidate identifier
skills	Candidate's technical and professional skills
qualification	Highest educational qualification
experience_level	Entry, Mid, Senior, or Expert experience category
job_role	Candidate's current or target role

3.3.3 LinkedIn Job Postings Dataset Features:

Feature	Description
job_id	Unique job-posting identifier
title	Free-text job title
description	Job responsibilities, requirements, and qualifications
company_name	Hiring company
location	Job location
formatted_experience_level	Required job experience category
skills_desc	Structured skill description where available
formatted_work_type	Employment type, such as full-time, contract, or internship
listed_time	Posting timestamp used for freshness
normalized_salary	Normalized salary where disclosed

The candidate dataset contains 98 unique skill combinations, 17 qualification categories, four experience-level categories, and 22 job-role categories. The job dataset contains 24,428 companies, 72,521 unique job titles, and 8,526 locations.

4. Data Governance

Area	Resume Parsing	Career Recommendation	Job Discovery and Matching
Data Source & Licensing	Publicly accessible Kaggle dataset published under the slug hadikp/resume-data-pdf . The exact licence and permitted use must be recorded from the Kaggle dataset page before submission	ESCO official public dataset	Candidate Job Role Dataset from Kaggle under MIT LinkedIn Job Postings Dataset from Kaggle under CC BY-SA 4.0
Ownership	Published by Kaggle user hadikp ; the underlying resumes appear to originate from resume-template and resume-builder sources	European Commission / ESCO classification	Candidate dataset published by C. K. Shetty; Job dataset published by Arsh Koneru and collaborators

Privacy	Resume images and annotations may contain names, emails, phone numbers, addresses, and profile links. These must be treated as potential PII	No personal data in ESCO	The candidate dataset contains structured profile attributes but no direct personal identifiers in the selected fields. Job descriptions may contain email addresses, URLs, or contact details, which are removed during text cleaning
Data Quality	The dataset contains 100% image-only PDFs, inconsistent folder labels, category imbalance, and perceptual near-duplicates	Missing values, duplicate records, unused skills, and invalid relations checked	Missing values, title inconsistencies, experience-level differences, sparse skill fields, duplicates, and text quality were analysed
Ethics & Bias	The data may contain template, language, geographic, occupational, and resume-format bias	ESCO may be Europe-focused and may not cover every local job title	Candidate data is synthetic and may not represent the full diversity of real applicants. Job data may reflect historical hiring, occupational, geographic, and seniority biases
Reproducibility	Dataset slug, download date, preprocessing notebook, random seed, canonical category mapping, perceptual-hash settings, schema, and split labels are documented	ESCO v1.2.1, English CSV classification used	Dataset names, versions, preprocessing notebooks, random seeds, retained columns, and generated files are documented

Privacy Measures

For resume-related data, sensitive information such as name, email, phone number, and address should be removed or masked before use. This is consistent with the project's Milestone 1 security plan, where personal information is masked before LLM processing.

Resume documents may contain personal contact details. Therefore, both the resume images and the extracted contact fields must be treated as sensitive data. Contact fields should not be included in public deliverables unless they have been masked or replaced with synthetic values.

Responsible Use

The system should recommend roles and rank jobs as decision-support only. It should not claim guaranteed hiring outcomes, fabricate skills, or use sensitive attributes for recommendation. The LinkedIn dataset is used for offline academic evaluation. The project does not use the dataset to bypass platform restrictions or automate access to protected pages. Any live job-discovery component will use only permitted public sources.

5. Exploratory Data Analysis

5.1 Resume Parsing Dataset EDA

EDA item	Result
Total resumes	8,905
Pages per resume	1 for all files
Raw category folders	97
Canonical categories	43
Corrupt files	0
Image-only PDFs	8,905 — 100%
Files with extractable text	0
Exact duplicate files	0
Perceptual groups	6,611
Files participating in multi-file near-duplicate groups	3,291
Excess copies beyond one representative per group	2,294
Largest category	Human Resources — 358
Smallest category	BPO — 29
Category imbalance ratio	Approximately 12.3:1
Median image resolution	962 × 1,600 pixels

Median file size	231 KB
Median OCR-estimated length	Approximately 374 words




Figure 1 shows file size (in KB) distribution in resume dataset

5.1.2 Category and Label Analysis

The original archive contains 97 folder names, but many represent different spelling or formatting variants of the same occupation category. Examples include “*Data Science*”, “*DataScience*”, and “*data science*” resumes. These labels were mapped to 43 canonical categories.

The category distribution is imbalanced. The largest category contains approximately 12.3 times as many resumes as the smallest category. To prevent majority categories from dominating evaluation, the gold set was sampled across all 43 canonical categories.

5.1.3 Duplicate Analysis

No exact byte-level duplicates were identified. However, perceptual hashing identified substantial visual duplication. The 8,905 files formed 6,611 perceptual groups. A total of 3,291 files belonged to groups containing more than one visually similar resume, representing 2,294 excess copies beyond one representative per group.

Perceptual groups were considered during gold-set selection to prevent visually similar documents from appearing in different evaluation splits.

5.1.4 Gold Annotation Field Coverage

Field group	Records containing the field
Contact information	84
Skills	86
Education	69
Experience	80
Projects	5
Certifications	27
Job titles	85

Projects and certifications are comparatively sparse. Therefore, final evaluation should report per-field results in addition to one overall F1-score.

5.1.5 Preliminary Common Skills

Based on case-insensitive exact matching in the current gold annotations, frequently occurring skills include:

- Communication — 11 resumes
- SQL — 9 resumes
- Leadership — 7 resumes
- Project Management — 7 resumes
- Python — 7 resumes
- SQL Server — 6 resumes
- Microsoft Office — 6 resumes
- AutoCAD — 6 resumes

This analysis is preliminary because synonyms such as MS Office and Microsoft Office have not yet been merged.

5.2 Career Recommendation Dataset EDA

Metric	Value
Clean occupations	3,039
Clean skills	13,939
Occupation-skill relation rows	126,051

Average skills per occupation	41.48
Minimum skills per occupation	1
Occupations with no linked skills	0
Unused skills	464

5.2.1 Missing Value Analysis

Missing values were handled during preprocessing. Empty alternative labels were replaced with blank strings. Occupation and skill descriptions were retained because they are useful for creating semantic embedding text.

5.2.2 Duplicate Analysis

Duplicate occupations and skills were removed using stable ESCO concept URIs. Duplicate occupation-skill relation pairs were also removed.

5.2.3 Feature Distribution

The main numerical feature is num_skills, which represents the number of skills linked with each occupation. Some occupations have many skills, while others have fewer. This distribution is useful because occupations with very broad skill coverage may dominate semantic matching if not handled carefully.




Figure 2 shows distribution of skills linked with each occupation

5.2.4 Common Skill Analysis

A skill frequency table was created to identify skills that appear across many occupations. Very common skills may be generic and less useful for distinguishing one occupation from another. Therefore, the preprocessing pipeline stores skill frequency and uses a specificity-based ordering while creating the embedding context.




Figure 3 shows top 20 most common skills in the ESCO dataset

5.3 Job Discovery and Matching Dataset EDA

5.3.1 Candidate Dataset

Metric	Value
Total candidate records	1,000
Total features	5
Unique skill combinations	98
Unique qualifications	17
Most common qualification	Bachelors in Computer Science — 349 candidates
Experience categories	4
Most common experience level	Mid — 401 candidates
Unique job roles	22
Most common job role	Mobile Developer — 75 candidates
Missing values	0
Exact duplicate rows	0

The qualification distribution is concentrated around Bachelor's degrees in Computer Science, Information Technology, and Software Engineering. Mid- and Senior-level

candidates are more common than Entry-level candidates. The job-role distribution is moderately balanced, although Mobile Developer is the most frequent class

5.3.2 LinkedIn Job Dataset Summary

Metric	Value
Total job postings	123,849
Original columns	31
Columns retained	14
Unique companies	24,428
Unique job titles	72,521
Unique locations	8,526
Missing skills_desc	121,410
Missing experience level	29,409
Missing salary fields	87,776
Missing company name	1,719
Missing description	7
Exact duplicate rows	0

The main data-quality issue is that skills_desc is missing for approximately 98% of job postings. Therefore, job title and full description are combined into a unified job_text field. Job titles also have extremely high cardinality, which shows that they are free-text values rather than standardized occupational labels. Salary values are right-skewed and contain high-value outliers, while experience-level and employment-type fields have lower cardinality and are more suitable for direct filtering or scoring.

6. Data Preprocessing

6.1 Resume Parsing Dataset Preprocessing

Step	Processing performed
Dataset indexing	Created one record for every PDF with its source folder, canonical category, filename, and file path
Category canonicalization	Converted 97 inconsistent folder names into 43 standard occupation categories
File validation	Checked that all PDFs could be opened and rendered
Exact duplicate detection	Generated file hashes to identify byte-identical documents

Perceptual duplicate detection	Rendered pages and generated perceptual hashes to group visually similar resumes
Image rendering	Rendered each PDF page using PyMuPDF at approximately 150–200 DPI
Image resizing	Downscaled very large images to control processing cost and memory usage
Image enhancement	Applied grayscale conversion and light denoising or contrast normalisation where required
OCR analysis	Used Tesseract on a sample for document-length and readability analysis; OCR was not used as the primary gold extraction method
Gold-set selection	Selected near-duplicate-safe representatives across all 43 canonical categories
Structured annotation	Stored gold profiles in JSON according to a fixed Candidate Profile schema
Output normalisation	Trimmed text, converted empty values to null, and deduplicated skills case-insensitively
Schema validation	Validated the profile object against <code>parsed_resume_schema.json</code>
Split assignment	Assigned gold records to development or held-out test sets using a fixed random seed

6.1.1 Model-specific Preparation

The Resume Parsing module uses a pretrained multimodal instruction-following model in a zero-shot, schema-constrained extraction setup. The model receives the resume image and returns structured JSON.

This module uses Gemma 3 as a module-specific multimodal parser, while the Career Recommendation and Job Matching modules may use different text or embedding models. The choice is justified because the Resume Parsing dataset is entirely image-based.

6.2 Career Recommendation Dataset Preprocessing

Step	Description
Column selection	Kept only useful columns such as occupation title, skill title, description, concept URI, relation type
Text cleaning	Removed extra spaces, line breaks, and inconsistent formatting
Missing value treatment	Filled missing alternative labels with blank strings
Duplicate removal	Removed duplicate occupations and skills using concept URI
Relation validation	Checked whether each occupation-skill relation points to a valid occupation and valid skill
Skill frequency calculation	Counted how often each skill appears across occupations

Specificity ordering	Ordered skills so more specific skills appear earlier in the embedding text
Feature engineering	Created num_skills, skills_list, and career_role_context
Final integration	Combined occupation, skill, and relation files into one final role-level dataset

6.3 Job Discovery and Matching Dataset Preprocessing

Step	Processing performed
Column selection	Removed columns not required for matching, ranking, traceability, or analysis
Missing-value handling	Filled missing experience and employment type with Unknown, company with Unknown Company, and salary with a sentinel value
Candidate skill parsing	Split comma-separated candidate skills, removed extra spaces, and created clean skill lists
Job text creation	Combined job title, skill description, and job description into a unified text field
Text cleaning	Converted text to lowercase, removed URLs, emails, special characters, generic stop words, and repeated whitespace
Duplicate removal	Removed duplicate job IDs, near-duplicate reposting, and repeated candidate profiles
Experience standardization	Mapped candidate and job experience levels to a shared scale
Job-title mapping	Used RapidFuzz to map free-text job titles to the 22 candidate role categories
Salary scaling	Applied Min-Max scaling to disclosed salary values using a scaler fitted only on training data
Categorical encoding	Used ordinal encoding for experience, one-hot encoding for employment type, and label encoding for mapped roles
Tokenization	Tokenized candidate skills and cleaned job text for BM25
Embedding generation	Generated 768-dimensional BGE embeddings for semantic similarity
Sparse-text filtering	Removed job records whose cleaned text remained too short for reliable matching

6.3.2 Shared Experience Scale

Raw job label	Standardized label
Internship	Entry
Entry level	Entry
Associate	Junior
Mid-Senior level	Mid
Director	Senior

Executive	Senior
Missing	Unknown

6.3.3 Engineered Ranking Signals

Signal	Description
bm25_score	Keyword relevance between candidate skills and job text
semantic	Cosine similarity between candidate and job embeddings
exp_fit	Compatibility of candidate and job experience levels
role_match	Fuzzy similarity between candidate target role and job title
freshness_score	Preference for recently posted jobs
final_score	Preliminary weighted combination of the ranking signals

7. Dataset Integration

Module	Integration Status
Resume Parsing	Dataset preparation completed; internal category and schema integration completed
Career Recommendation	Completed for ESCO files
Job Discovery and Matching	Completed

a. Resume parsing:

No external dataset was merged into the Resume Parsing corpus. However, internal integration was required because the source archive contained multiple naming conventions and visually duplicated resume templates.

The integration process included:

1. mapping 97 raw folder names to 43 canonical categories;
2. linking visually similar resumes through a shared dup_group;
3. converting all gold annotations into one common Candidate Profile schema;

b. Career Recommendation: three main ESCO files were integrated:

1. occupations_en.csv
2. skills_en.csv
3. occupationSkillRelations_en.csv

These were joined using ESCO occupation URIs and skill URIs.

c. Job Matching: The candidate and job datasets are not merged into one ordinary flat table. They are integrated using a query-document structure

- candidate profile = query,
- job posting = document,
- candidate-job score = relevance estimate.

Schema alignment was performed as follows:

Candidate field	Job field	Alignment
experience_level	formatted_experience_level	Mapped to shared experience categories
job_role	title	Fuzzy title-to-role matching
skills	skills_desc + description	Converted into normalized text and tokens
Candidate text	Job text	Converted into embeddings for cosine similarity

Because skills_desc is mostly missing, the job description acts as the primary fallback source of skill and requirement information

8. Data Augmentation

8.1 Resume Parsing Module

No image augmentation was applied to the gold development or test sets because evaluation must reflect the original documents.

Controlled augmentation may be considered later for robustness testing, including:

- slight image rotation;
- blur;
- JPEG compression;
- contrast reduction;
- resolution reduction.

Any future augmented version must remain outside the held-out test set and must retain a reference to its original document to prevent leakage.

8.2 Job Discovery and matching module

Positive training pairs are augmented using controlled skill-synonym expansion. Selected skill terms are replaced with equivalent expressions, such as:

Skill	Variants
Python	Python programming, Python scripting, Python development
SQL	Structured Query Language, SQL queries, database queries
Machine Learning	ML, supervised learning, predictive modelling
Power BI	PowerBI, business intelligence, BI dashboards
Tableau	Tableau Desktop, Tableau visualisation, data visualisation

Two augmented versions are generated for each positive training pair. Augmentation is applied only to the training split. Validation and test data remain unchanged.

9. Dataset Splitting

9.1 Resume Parsing

The Resume Parsing module uses a pretrained model without fine-tuning. Therefore, no gradient-based training split is required. The prepared splits are used for prompt development and final evaluation.

Dataset portion	Size	Purpose
Training set	0	No model weights are trained
Development set	34	Prompt development, output-format checks, and error analysis
Test set	52	Final held-out Precision, Recall, and F1 evaluation
Unannotated source pool	8,819	EDA, future sampling, and robustness experiments
Total gold set	86	Structured evaluation records covering 43 categories

9.2 Career Recommendation

The ESCO occupation dataset itself is not split because it acts as the complete reference knowledge base for retrieval. Splitting will be applied to candidate profiles during evaluation.

9.3 Job Discovery and Matching

In candidate dataset, Candidate IDs are divided before candidate-job pair construction to prevent the same candidate from appearing in multiple splits.

Split	Candidate ratio	Candidate count from 200	Purpose
Development/Train	60%	120	Develop and tune ranking logic
Validation	20%	40	Select weights and thresholds
Test	20%	40	Final unbiased evaluation

10. Challenges Encountered

Challenge	Module	Handling
Resume data may contain PII	Resume Parsing	To be handled by anonymization and masking
Resume formats vary widely	Resume Parsing	To be handled by text extraction and layout-aware checks

No direct candidate-to-ESCO public ground truth	Career Recommendation	Manual candidate-to-ESCO mapping will be prepared
Duplicate and unused ESCO entries	Career Recommendation	Duplicate removal and unused skill reporting completed
Some skills are too generic	Career Recommendation	Skill frequency and specificity ordering added
ESCO is Europe-focused	Career Recommendation	Accepted as a standard taxonomy, but limitation documented
skills_desc missing in about 98% of jobs	Job Matching	Created job_text from title and description
Missing job experience level	Job Matching	Mapped to Unknown and applied neutral scoring
Negative-positive imbalance	Job Matching	Applied training-only positive-pair augmentation
Salary missing in most postings	Job Matching	Treated as optional information rather than a hard filter
Long job descriptions	Job Matching	Limited embedding input to the most relevant text within the model context limit
Potential low-confidence title mapping	Job Matching	Retained mapping confidence and will analyse low-confidence cases
Image-only coverage	Resume Parsing	Native-text PDF and DOCX resumes remain a future evaluation requirement
Gold-schema mismatch	Resume Parsing	Validate only the nested profile object or introduce a separate gold-record schema

11. Deliverable Produced

11.1 Resumes Dataset – [Click Here](#)

11.2 Career Recommendation Dataset – [Click Here](#)

11.3 Job Matching Dataset – [Click Here](#)

12. Summary and Next Steps

Milestone 2 established the data foundation for the three main project modules. The Career Recommendation dataset was prepared from ESCO occupations, skills, and occupation-skill relations. The Job Discovery and Matching datasets were prepared using structured candidate profiles and LinkedIn job postings. EDA identified missing skill descriptions, unstandardized job titles, inconsistent experience labels, missing salary data, and class imbalance. These issues were addressed through text construction, normalization, fuzzy role mapping,

controlled missing-value handling, deduplication, feature engineering, and candidate-job pair construction.

The Career Recommendation data is ready for ESCO occupation embedding and ChromaDB retrieval. The Job Matching data is ready for preliminary BM25, embedding, and hybrid-ranking experiments. Proxy relevance labels will be used only for initial development, while a manually labelled evaluation subset will be prepared for final Milestone 3 evaluation.

13. Team Consent

Name	Roll Number	Initial
Gaurav Kumar	22f1001105	GK
Dev Gupta	22f2000888	DG
Abhinav Ohri	24f1002064	AO
Pranav N	22f2000117	PN
Praveena N	22f3001454	PN